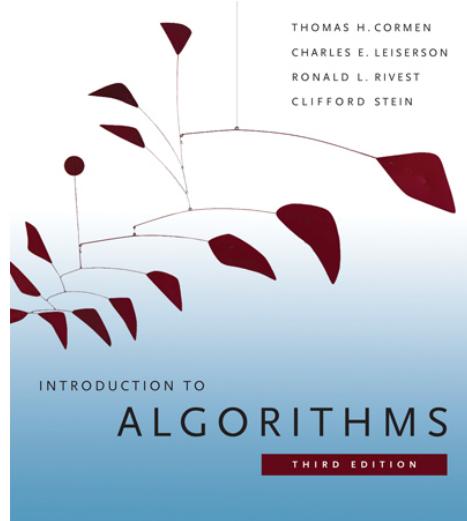


6.006- *Introduction to Algorithms*



Lecture 6

Prof. Manolis Kellis
CLRS: Chapter 17 and 32.2.

Unit #2 – Genomes, Hashing, and Dictionaries

Unit	Pset	Week	Date	Lecture (Tuesdays and Thursdays)		Recitation (Wed and Fri)	
Intro Binary Search Trees	PS1 Out: 2/1 Due: Mon 2/14 HW lab: Sun 2/13	1	Tue Feb 01	1	Introduction and Document Distance	1	Python and Asymptotic Complexity
			Thu Feb 03	2	Peak Finding Problem	2	Peak Finding correctness & analysis
		2	Tue Feb 08	3	Scheduling and Binary Search Trees	3	Binary Search Tree Operations
			Thu Feb 10	4	Balanced Binary Search Trees	4	Rotations and AVL tree deletions
Hashing	PS2 Out: 2/15 Due: Mon 2/28 HW lab: Sun 2/27	3	Tue Feb 15	5	Hashing I : Chaining, Hash Functions	5	Hash recipes, collisions, Python dicts
			Thu Feb 17	6	Hashing II : Table Doubling, Rolling Hash	6	Probability review, Pattern matching
		4	Tue Feb 22	-	President's Day - Monday Schedule - No Class	-	No recitation
			Thu Feb 24	7	Hashing III : Open Addressing	7	Universal Hashing, Perfect Hashing
Sorting	PS3. Out: 3/1 Due: Mon 3/7 HW lab: Sun 3/6	5	Tue Mar 01	8	Sorting I : Insertion & Merge Sort, Master Theorem	8	Proof of Master Theorem, Examples
			Thu Mar 03	9	Sorting II : Heaps	9	Heap Operations
		6	Tue Mar 08	10	Sorting III: Lower Bounds, Counting Sort, Radix Sort	10	Models of computation
			Wed Mar 09	Q1	Quiz 1 in class at 7:30pm. Covers L1-R10. Review Session on Tue 3/8 at 7:30pm.		
Graphs and Search	PS4. Out: 3/10 Due: Fri 3/18 HW lab: W 3/16	7	Thu Mar 10	11	Searching I: Graph Representation, Depth-1st Search	11	Strongly connected components
			Tue Mar 15	12	Searching II: Breadth-1st Search, Topological Sort	12	Rubik's Cube Solving
		9	Thu Mar 17	13	Searching III: Games, Network properties, Motifs	13	Subgraph isomorphism
			Tue Apr 05	14	Shortest Paths I: Introduction, Bellman-Ford	14	Relaxation algorithms
Shortest Paths	PS5 Out: 3/29 Due: Mon 4/11 HW lab: Sun 4/10	8	Thu Mar 31	15	Shortest Paths II: Bellman-Ford, DAGs	15	Shortest Path applications
			Tue Apr 07	16	Shortest Paths III: Dijkstra	16	Speeding up Dijkstra's algorithm
		10	Thu Apr 07	17	Graph applications, Genome Assembly	17	Euler Tours
			Tue Apr 12	18	DP I: Memoization, Fibonacci, Crazy Eights	18	Limits of dynamic programming
Dynamic Program ming	PS6 Out: Tue 4/12 Due: Fri 4/29 HW lab: W 4/27	11	Wed Apr 13	Q2	Quiz 2 in class at 7:30pm. Covers L11-R17. Review Session on Tue 4/13 at 7:30pm.		
			Thu Apr 14	19	DP II: Shortest Paths, Genome sequence alignment	19	Edit Distance, LCS, cost functions
		12	Tue Apr 19	-	Patriot's Day - Monday and Tuesday Off	-	No recitation
			Thu Apr 21	20	DP III: Text Justification, Knapsack	20	Saving Princess Peach
Numbers Pictures (NP)	PS7 out Thu 4/28 Due: Fri 5/6 HW lab: Wed 5/4	13	Tue Apr 26	21	DP IV: Piano Fingering, Vertex Cover, Structured DP	21	Phylogeny
			Thu Apr 28	22	Numerics I - Computing on large numbers	22	Models of computation return!
		14	Tue May 3	23	Numerics II - Iterative algorithms, Newton's method	23	Computing the nth digit of π
			Thu May 5	24	Geometry: Line sweep, Convex Hull	24	Closest pair
Beyond		15	Tue May 10	25	Complexity classes, and reductions	25	Undecidability of Life
			Thu May 12	26	Research Directions (15 mins each) + related classes		
			Finals week	Q3	Final exam is cumulative L1-L26. Emphasis on L18-L26. Review Session on Fri 5/13 at 3pm		

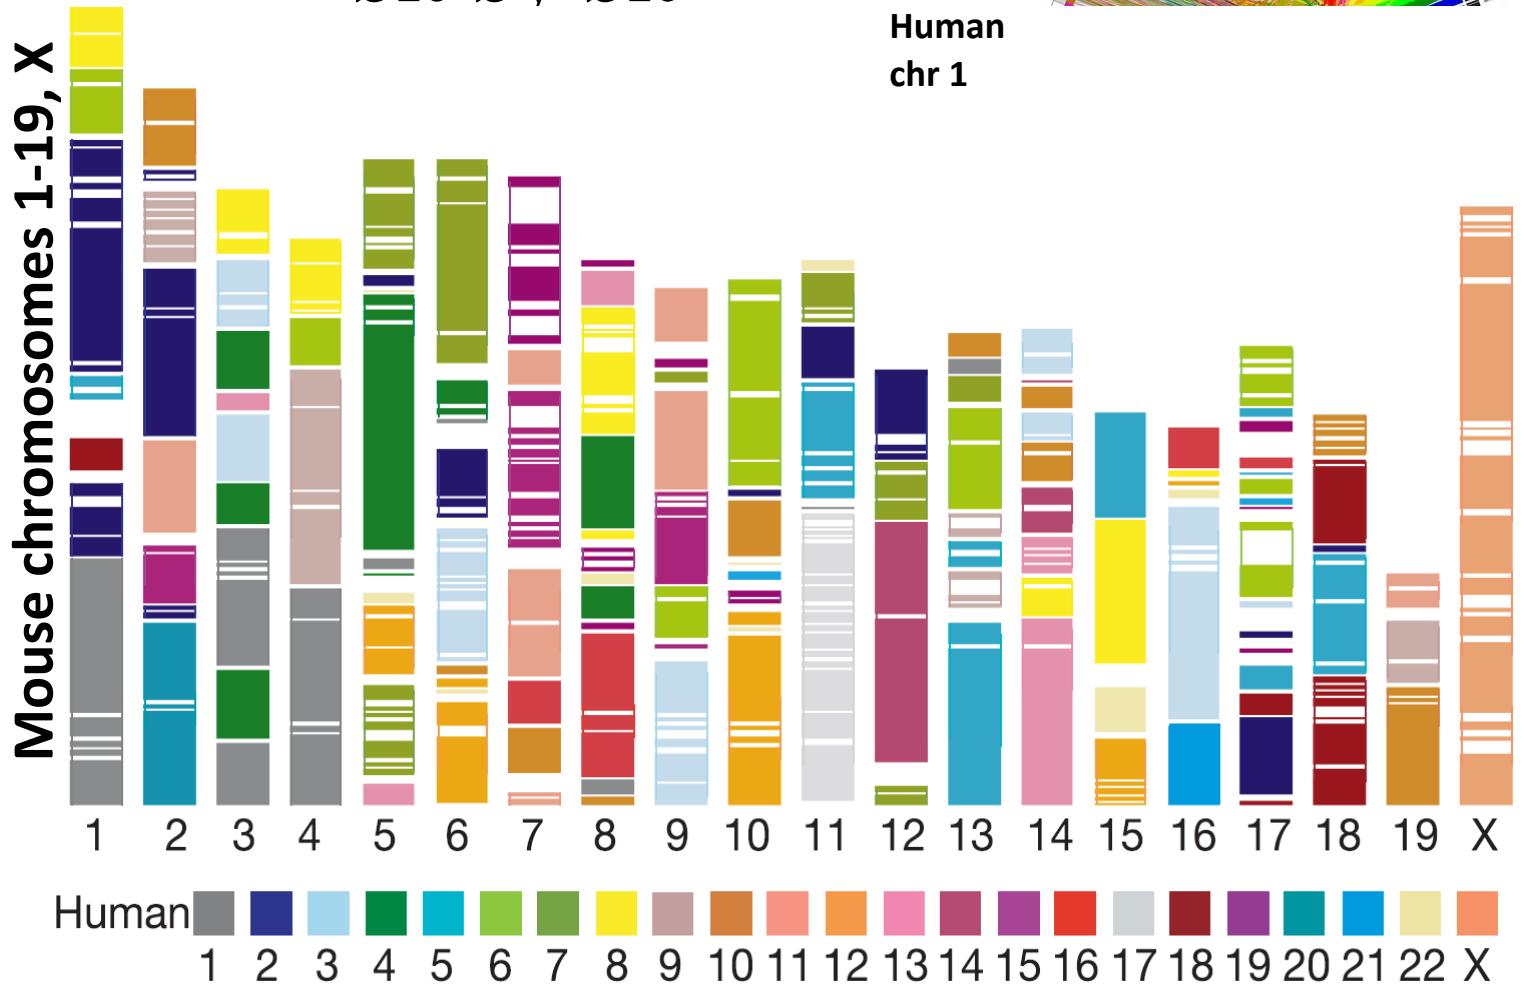
Unit #2: Hashing

- **Last time: Genomes, Dictionaries, Hashing**
 - Intro, basic operations, collisions and chaining
 - Simple uniform hashing assumption
- **Today: Faster hashing, hash functions**
 - Hash functions, python implementation
 - Faster comparison: Signatures
 - Faster hashing: Rolling Hash
- **Next week: Space issues**
 - Dynamic resizing and amortized analysis
 - Open addressing, deletions, and probing

Our plan for today: Hashing II

- **Today: Hashing in practice,**
 - Review: Genomes, Dictionaries, Hashing
 - Review: Hash functions, collisions, chaining
 - Hash functions in practice: div, mult, Python
 - Speeding up hashing: $O(n^4) \rightarrow O(n^2 \lg n) \rightarrow O(n)$
 - Faster hashing: rolling hash
 - Fewer string comparisons: signatures
 - Summary, conclusion
- **Next week: Space issues**

Comparing two genomes bit by bit



Naïve Algorithm

- Say strings S and T of length n
- For $L = n$ downto 1
 - for all length L substrings X_1 of S n^*
 - for all length L substrings X_2 of T n^*
 - if $X_1 = X_2$, return L n
- Runtime analysis
 - n candidate lengths
 - n strings of that length in X_1
 - n strings of that length in X_2
 - L time to compare the strings
 - Total runtime: $\mathcal{O}(n^4)$

Dictionaries: Formal Definition

- It is a **set** containing **items**; each item has a **key**
- what keys and items are is quite flexible
- Supported Operations:
 - **Insert(*key*, *item*)**: add *item* to set, indexed by *key*
 - **Delete(*key*)**: delete item indexed by *key*
 - **Search(*key*)**: return the item corresponding to the given *key*, if such an item exists
 - **Random_key()**: return a random key in dictionary
- **Assumption**: every item has its own key (or that inserting new item clobbers old)
- **Application** (and origin of name): Dictionaries
 - *Key* is word in English, *item* is word in French

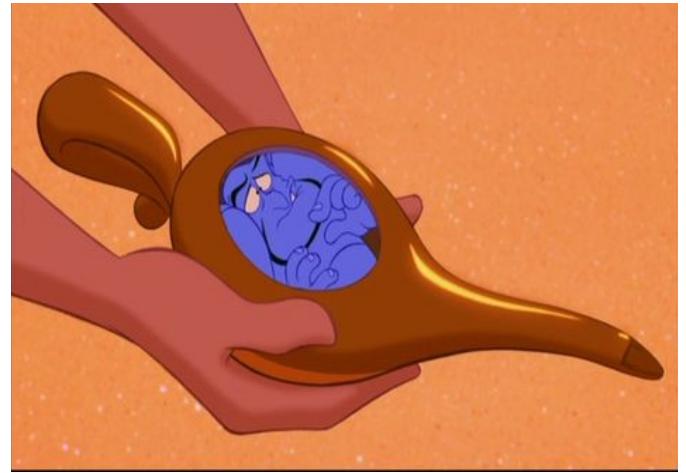
Interpreting words as numbers

- What if keys aren't numbers?
 - Anything in the computer is a sequence of bits
 - So we can pretend it's a number
- Example: English words
 - 26 letters in alphabet
 - ⇒ can represent each with 5 bits
 - Antidisestablishmentarianism has 28 letters
 - $28 \times 5 = 140$ bits
 - So, store in array of size 2^{140} oops
- **Why hashing:** 2^{140} is too much space when in fact only 100,000 words actually appear

Our plan for today: Hashing II

- **Today: Hashing in practice,**
 - Review: Genomes, Dictionaries, Hashing
 - Review: Hash functions, collisions, chaining
 - Hash functions in practice: div, mult, Python
 - Speeding up hashing: $O(n^4) \rightarrow O(n^2 \lg n) \rightarrow O(n)$
 - Faster hashing: rolling hash
 - Fewer string comparisons: signatures
 - Summary, conclusion
- **Next week: Space issues**

The ‘magic’ of hash functions



PHENOMENAL
COSMIC
POWERS!!

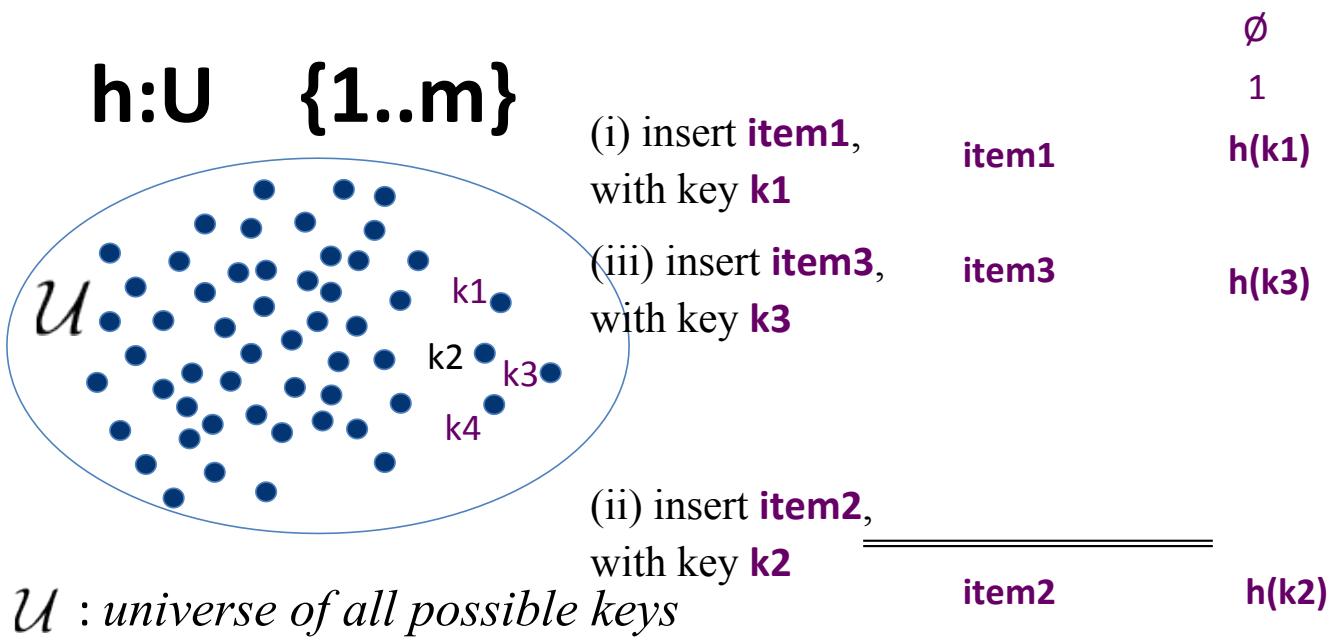


itty bitty living space

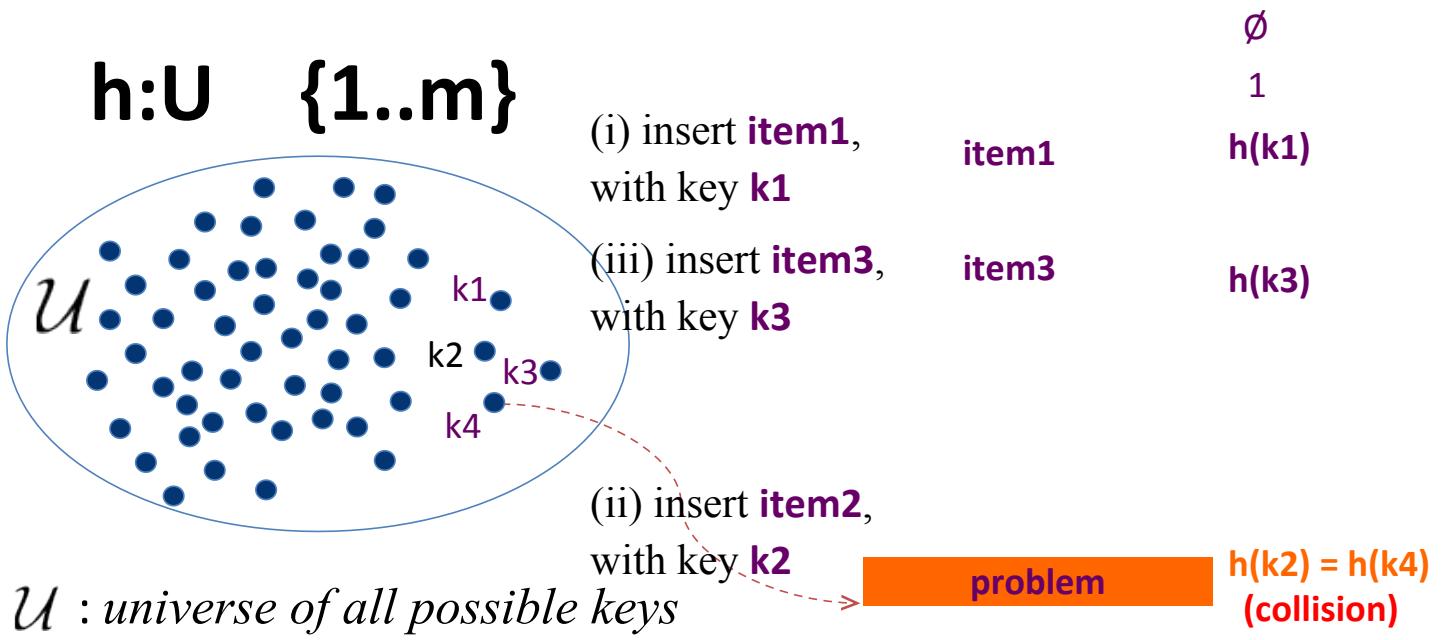
With apologies to Disney

Huge universe of possible keys...

... all map to small table m



Huge universe of possible keys ... all map to small table m ... leading to collisions

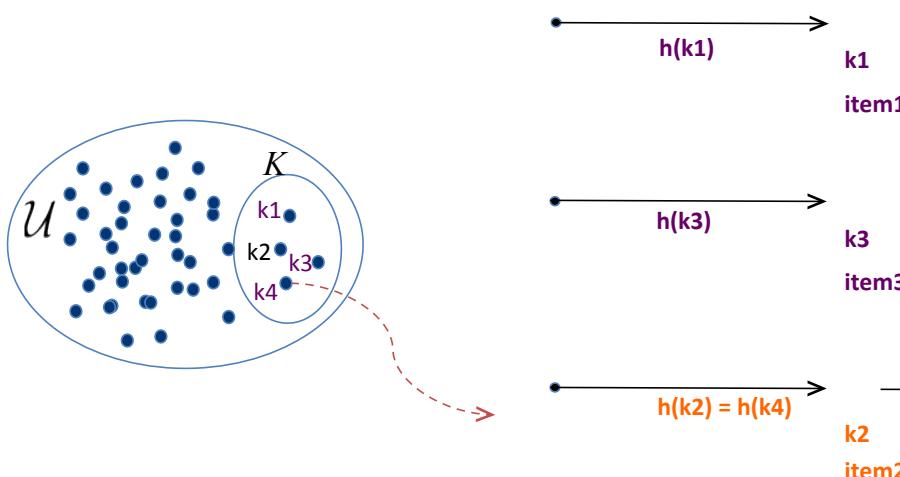


m-1

Coping with collisions: Chaining

\mathcal{U} : universe of all possible keys

K : actual keys, not known in advance



- Each bucket, linked list of contained items
- Space used is *space of table plus one unit per item (size of key and item)*

How long will 1 lookup take? SUHA
 - Worst case: **O(n)** Average Case: **O(1)?**

E.g. During look-up of **k4**:
 1. Hash key **k4**, obtaining address **h(k4)**, traverse linked list
 2. Traverse linked list, **compare k4 to k2**: mismatch → continue
 3. **Compare k4 to k4**: match → return **item4**

Simple uniform hashing assumption

- Each key $k \in K$ of keys is equally likely to be hashed to any slot of table T , independent of where other keys are hashed.

Let n be the number of keys in the table, and let m be the number of slots.

Define the *load factor* of T to be $\alpha = n/m$

= average number of keys per slot.

Expected lookup time in average case is $E(..) = O(1 + \alpha)$

If $m = \Omega(n) \rightarrow \alpha = O(1) \rightarrow E(..) = O(1)$

*apply hash function
and access the slot*

*search the
linked list*

Our plan for today: Hashing II

- **Today: Hashing in practice,**
 - Review: Genomes, Dictionaries, Hashing
 - Review: Hash functions, collisions, chaining
 - Hash functions in practice: div, mult, Python
 - Speeding up hashing: $O(n^4) \rightarrow O(n^2 \lg n) \rightarrow O(n)$
 - Faster hashing: rolling hash
 - Fewer string comparisons: signatures
 - Summary, conclusion
- **Next week: Space issues**

Problem: Reality is not always nice

- Keys are often very nonrandom
 - Regularity (evenly spaced sequence of keys)
 - All sorts of mysterious patterns
- Solution: pick a hash function whose values “look” random
 - Though a given key must always have same hash
- Similar to pseudorandom generators
 - Answer based on key & random seed for whole table
- For any function, some set of keys that is bad
 - but hopefully not your set
- In practice, this is all about good *heuristics!*

Division Hash Function

- $h(k) = k \bmod m$
- k_1 and k_2 collide when $k_1 = k_2 \pmod{m}$
 - Unlikely if keys are random
- Problem: if m is a power of 2, the hash doesn't even depend on all the bits of k

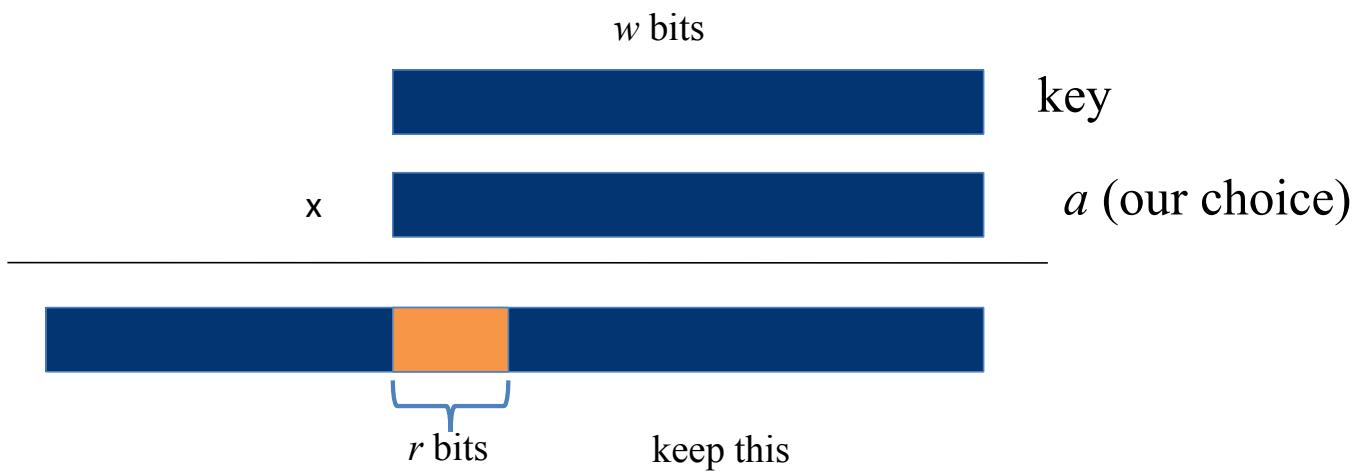
If $k = 1011000111\textcolor{blue}{011010}_2$ and $r = 6$, then
 $h(k) = 011010_2$. $\textcolor{red}{h(k)}$

Problems with division method

- Regularity
 - Suppose keys are $x, 2x, 3x, 4x, \dots$
 - Suppose x and chosen m have common divisor d
 - Then only use $1/d$ fraction of table
 - x series cycles back, leaving $d-1$ out of d entries blank
 - E.g, m power of 2 and all keys are even, only use half
- So make m a prime number
 - But finding a prime number is hard
 - And now you have to divide (slow)

Multiplication Hash Function

- Suppose we're aiming for table size 2^r
- and keys are w bits long, where $w > r$ is the machine word
- Multiply k with some a (fixed for the hash function)
- then keep certain bits of the result as follows

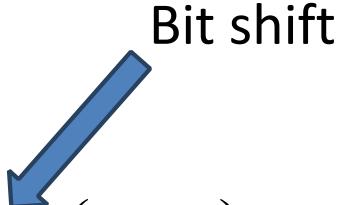


Multiplication Hash Function

- The formula:

$$h(k) = [(a * k) \bmod 2^w] \gg (w - r)$$

- Multiply by a
- When overflow machine word, wrap
- Take high r bits of resulting machine word
- (Assumes table size smaller than machine word)

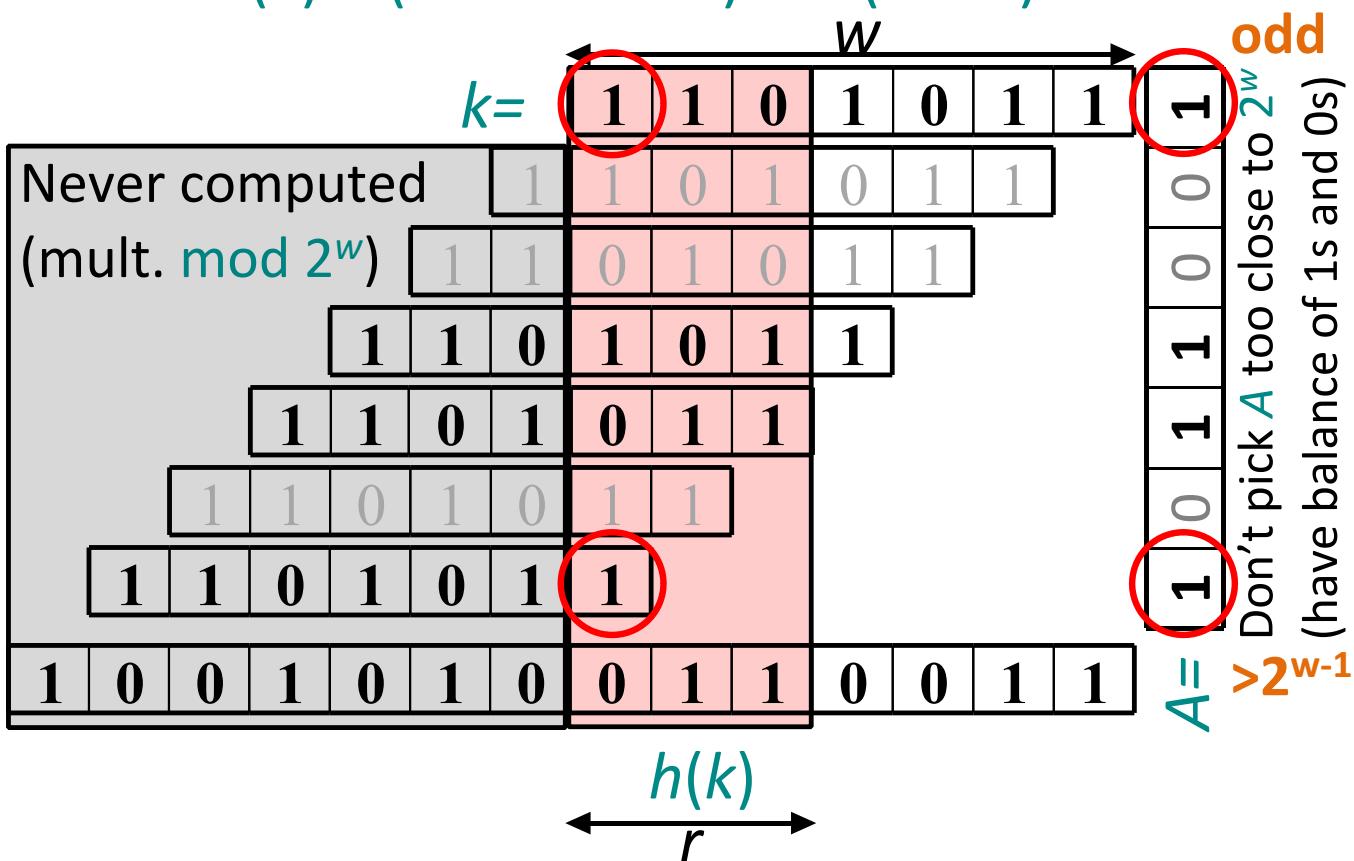


Benefit: Multiplying and bit shifts faster than division

Good practice: Make a an odd integer (why?) $> 2^{w-1}$

Multiplication method intuition (1)

$$h(k) = (A \cdot k \bmod 2^w) \text{ rsh } (w - r)$$

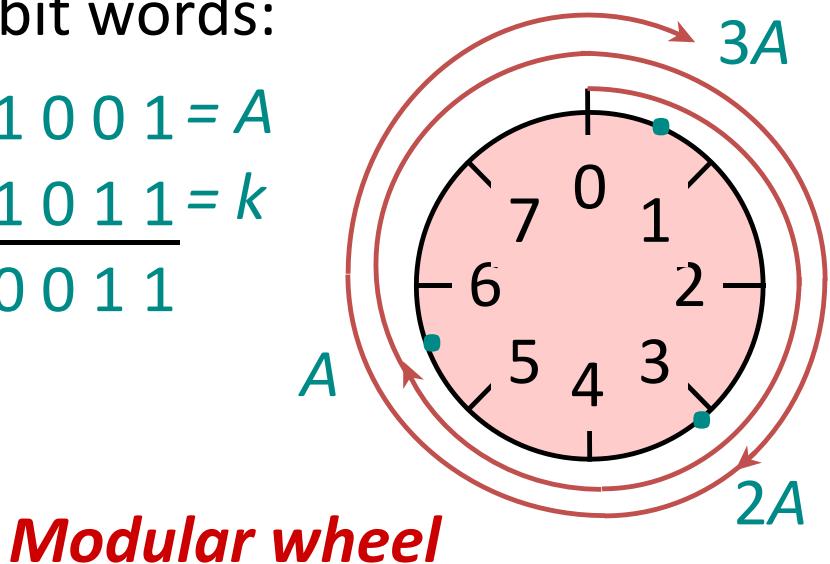


Multiplication method intuition (2)

$$h(k) = (A \cdot k \bmod 2^w) \text{ rsh } (w - r)$$

Suppose that $m = 8 = 2^3$ and that our computer has $w = 7$ -bit words:

$$\begin{array}{r} 1011001 = A \\ \times \quad 1101011 = k \\ \hline 10010100110011 \\ h(k) \end{array}$$



Don't pick A too close to 2^w
(explore more of the wheel space)

Python Implementation

- Python objects have a hash method
 - Number, string, tuple, any object implementing `__hash__`
- Maps object to (arbitrarily large) integer
 - So really, should be called prehash
- Take mod m to put in a size-m hash table
- Peculiar details
 - Integers map to themselves
 - Strings that differ by one letter don't collide
 - “better” than random for common sequences
 - 1,2,3,4,5 or var_a, var_b, var_c, var_d

Our plan for today: Hashing II

- **Today: Hashing in practice,**
 - Review: Genomes, Dictionaries, Hashing
 - Review: Hash functions, collisions, chaining
 - Hash functions in practice: div, mult, Python
 - Speeding up hashing: $O(n^4) \rightarrow O(n^2 \lg n) \rightarrow O(n)$
 - Faster hashing: rolling hash
 - Fewer string comparisons: signatures
 - Summary, conclusion
- **Next week: Space issues**

Recall: Longest Common Substring

- Strings S, T of length n, want to find longest common substring
- Algorithms from last time:
 - $O(n^4) \rightarrow O(n^3 \log n) \rightarrow O(n^2 \log n)$
 - Naïve \rightarrow Binary Search on L \rightarrow Hash table for T
- Winner algorithm used a hash table of size n:
Binary search on maximum match length L; to check if a length works:
 - Insert all length-L substrings of S in hash table
 - For each length-L substring x of T
 - Look in bucket $h(x)$ to see if x is in S

Analysis?

Runtime Analysis: bottlenecks

- Binary search cost: $O(\log n)$ length values L tested
- For each length value L , here are the costly operations:
 - Inserting all L -length substrings of S : $n-L$ hashes
 - Each hash takes L time, so total work $Q((n-L)L)=Q(n^2)$
 - Hashing all L -length substrings of T : $n-L$ hashes
 - another $Q(n^2)$
 - Time for comparing substrings of T to substrings of S :
 - How many comparisons?
 - Under SUHA, each substring of T is compared to an expected $O(1)$ of substrings of S found in its bucket
 - Each comparison takes $O(L)$
 - Hence, time for all comparisons: $Q(nL)=Q(n^2)$
- So $Q(n^2)$ work for each length
- Hence $\Omega(n^2 \log n)$ including binary search

Making comparisons even faster?

- Amdahl's law: if one part of the code takes 20% of the time, then no matter how much you improve it, you only get 20% speedup
- Corollary: must improve **all** asymptotically worst parts to change asymptotic runtime
- In our case
 - Must compute sequence of n hashes faster: $n^2 \rightarrow n$
 - Must reduce cost of comparing in bucket: $n^2 \rightarrow n$

Our plan for today: Hashing II

- **Today: Hashing in practice,**
 - Review: Genomes, Dictionaries, Hashing
 - Review: Hash functions, collisions, chaining
 - Hash functions in practice: div, mult, Python
 - Speeding up hashing: $O(n^4) \rightarrow O(n^2 \lg n) \rightarrow O(n)$
 - Faster hashing: rolling hash
 - Fewer string comparisons: signatures
 - Summary, conclusion
- **Next week: Space issues**

Rolling Hash

- We make a sequence of n substring hashes
 - Substring lengths L
 - Total time $O(nL) = O(n^2)$
- Can we do better?
 - For our particular application, yes!

length n

Verba volant, scripta manent

length L

Rolling Hash Idea

- e.g. hash all 3-substrings of “there”
- Recall division hash: $x \bmod m$
- Recall string to number:
 - First substring “the” = $t \cdot (26)^2 + h \cdot (26) + e$
- If we have “the”, can we compute “her” as $f(\text{“the”})$?

$$\begin{aligned}\text{“her”} &= h \cdot (26)^2 + e \cdot (26) + r \\ &= 26 \cdot (h \cdot (26) + e) + r \\ &= 26 \cdot (t \cdot (26)^2 + h \cdot (26) + e - t \cdot (26)^2) + r \\ &= 26 \cdot (\text{“the”} - t \cdot (26)^2) + r\end{aligned}$$

- i.e. subtract first letter’s contribution to number, shift, and add last letter

General rule for $s[i+1..] = f(s[i..])$

Start:

$S[i]$ $S[i+1]$ $S[i+2]$... $S[i+L-2]$ $S[i+L-1]$

1. Subtract 1st

Finish:

$S[i+1]$ $S[i+2]$... $S[i+L-2]$ $S[i+L-1]$ $S[i+L]$

2. Shift

3. Add last

- Strings = base- b numbers
- Start: Current substring $S[i \dots i+L-1]$

$$S[i] \cdot b^{L-1} + S[i+1] \cdot b^{L-2} + S[i+2] \cdot b^{L-3} \dots + S[i+L-1]$$

1 - $S[i] \cdot b^{L-1}$ *subtract higher-order term*

$$S[i+1] \cdot b^{L-2} + S[i+2] \cdot b^{L-3} \dots + S[i+L-1]$$

2 * b *shift to the left*

$$S[i+1] \cdot b^{L-1} + S[i+2] \cdot b^{L-2} \dots + S[i+L-1] \cdot b$$

3 + $S[i+L]$ *add lower-order term*

$$S[i+1] \cdot b^{L-1} + S[i+2] \cdot b^{L-2} \dots + S[i+L-1] \cdot b + S[i+L]$$

- Finish: Next substring $S[i+1 \dots i+L]$

Ensuring O(1) cut/shift/paste operations

Start:

$S[i] \ S[i+1] \ S[i+2] \ \dots \ S[i+L-2] \ S[i+L-1]$

1. Subtract 1st
2. Shift
3. Add last

Finish:

$S[i+1] \ S[i+2] \ \dots \ S[i+L-2] \ S[i+L-1] \ S[i+L]$

- So: $S[i+1 \ \dots \ i+L] = b \ S[i \ \dots \ i+L-1] - b^L \ S[i] + S[i+L]$
- where
- $S[i \ \dots \ i+L-1] = S[i] \cdot b^{L-1} + S[i+1] \cdot b^{L-2} + \dots + S[i+L-1] \ (*)$
- **But** $S[i \ \dots \ i+L-1]$ may be a huge number (so huge that we may not even be able to store in the computer, e.g. $L=50$, $b=26$)
- **Solution** only keep its *division hash*: $S[\dots] \bmod m$
- This can be computed without computing $S[\dots]$, using **mod magic!**
- Recall: $(ab) \bmod m = (a \bmod m) (b \bmod m) \bmod m$
 $(a+b) \bmod m = ((a \bmod m) + (b \bmod m)) \bmod m$
- With a clever parenthesization of (*): O(L) to hash string!

Applying mod magic in hash space

Start:

$S[i] \ S[i+1] \ S[i+2] \ \dots \ S[i+L-2] \ S[i+L-1]$

1. Subtract 1st

Finish:

$S[i+1] \ S[i+2] \ \dots \ S[i+L-2] \ S[i+L-1] \ S[i+L]$

2. Shift

3. Add last

- Recall: $S[i+1 \dots i+L] = b S[i \dots i+L-1] - b^L S[i] + S[i+L]$
- But can we also compute:
 $\text{hash}(S[i+1 \dots i+L])$ from $\text{hash}(S[i \dots i+L-1])$?
- Still mod magic to the rescue!
- $$h(S[i+1 \dots i+L]) = h(b S[i \dots i+L-1]) - h(b^L \cdot S[i]) + h(S[i+L])$$

Can compute $b^L \bmod m$ in $O(L)$ by recursive multiplication mod m

Total time: $O(L)$ time for the first hash
 $+O(L)$ to compute $b^L \bmod m$
 $+L*O(1)$ for each additional hash



Computing $n-L$ hashes costs $O(n)$

Our plan for today: Hashing II

- **Today: Hashing in practice,**
 - Review: Genomes, Dictionaries, Hashing
 - Review: Hash functions, collisions, chaining
 - Hash functions in practice: div, mult, Python
 - Speeding up hashing: $O(n^4) \rightarrow O(n^2 \lg n) \rightarrow O(n)$
 - Faster hashing: rolling hash
 - Fewer string comparisons: signatures
 - Summary, conclusion
- **Next week: Space issues**

Faster Comparison

- **First Idea:** when we find a match for some length, we can stop and go to the next value of length in our binary search.
- **But,** the real problem is “false positives”
 - Strings in same bucket that don’t match, but we waste time on
- **Analysis:**
 - n substrings to size- n table: average load **1**
 - SUHA: for every substring x of T , there is **1** other string in x ’s bucket (in expectation)
 - Comparison work: **L** per string (in expectation)
 - So total work for all strings of T : **$nL = Q(n^2)$**

Solution: Bigger table!

- What size?
- Table size $\mathbf{m = n^2}$
 - n substrings to size- m table: average load $\mathbf{1/n}$
 - SUHA: for every substring x of T , there is $\mathbf{1/n}$ other strings in x 's bucket (in expectation)
 - Comparison work: $\mathbf{L/n}$ per string (in expectation)
 - So total work for all strings of T : $\mathbf{n(L/n) = L = O(n)}$
- Downside?
 - Bigger table
 - (n^2 isn't realistic for large n)

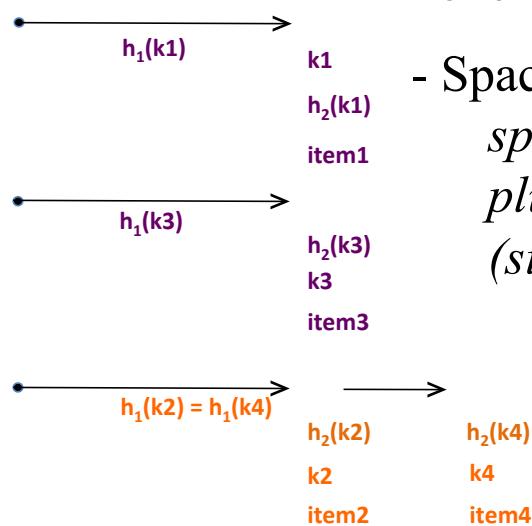
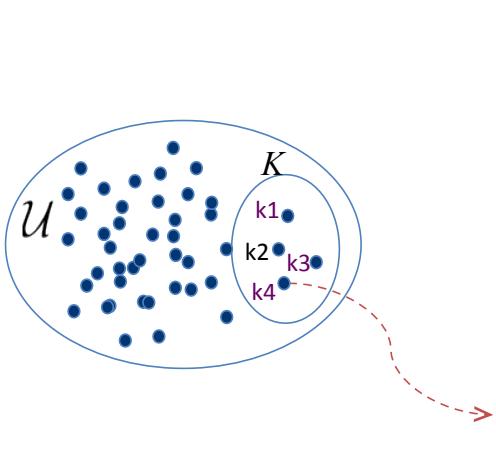
Signatures

- Note n^2 table isn't needed for fast lookup
 - Size n enough for that
 - n^2 is to reduce cost of false positive compares
- So don't bother making the n^2 table
 - Just compute for each string another hash **value** in the larger range $1..n^2$
 - Called a **signature**
 - If two signatures differ, strings differ
 - $\Pr[\text{same sig for two different strings}] = 1/n^2$
 - (simple uniform hashing)

Chaining with signatures

\mathcal{U} : universe of all possible keys

K : actual keys, not known in advance



- Each bucket, linked list of contained items
- Space used is
space of table plus one unit per item (size of key and item)

How long will 1 lookup take?

- Worst case: $O(n)$ Average Case: $O(1)$? SUHA

E.g. During look-up of **k4**:

1. Hash key **k4** with h_1 and h_2 obtaining address and signature
2. Traverse linked list, **compare signatures**: mismatch \rightarrow cont.
3. **Compare k4 to k4**: match \rightarrow return **item4**

Application

- Hash substrings to size n table
- But store a signature with each substring
 - Using a second hash function to $[1..n^2]$
- Check each T-string against its bucket
 - First check signature, if match then compare strings
 - Signature is a small number, so comparing them is $O(1)$

strictly speaking $O(\log n)$; but if $n^2 < 2^{32}$ the signature fits inside a word of the computer; in this case, the comparison takes $O(1)$

Application

- Runtime Analysis:
 - for each T-string:
 $O(\text{bucket size})=O(1)$ work to compare **signatures**;
 - so overall $O(n)$ time in signature comparisons
 - Time spent in string comparisons?
 $L \times (\text{Expected Total Number of False-Signature Collisions})$
 - n out of the n^2 values in $[1..n^2]$ are used by S-strings
 - so probability of a T-string signature-colliding with some S-string: n/n^2
 - hence total expected number of collisions 1
- so total time spent in String Comparisons is L
- fine print:** time needed to compute signatures using rolling hash $O(n)$ for each signature, hence total time $O(2n)=O(n)$. (again, idea to optimize bottlenecks, by paying small incremental costs in non-bottleneck spots)

Our plan for today: Hashing II

- **Today: Hashing in practice,**
 - Review: Genomes, Dictionaries, Hashing
 - Review: Hash functions, collisions, chaining
 - Hash functions in practice: div, mult, Python
 - Speeding up hashing: $O(n^4) \rightarrow O(n^2 \lg n) \rightarrow O(n)$
 - Faster hashing: rolling hash
 - Fewer string comparisons: signatures
 - Summary, conclusion
- **Next week: Space issues.** Dynamic resizing, open addressing, probe sequences, deletions

Summary (rehash)

- Matching big genomes is a hard problem
 - And you will tackle it in your problem set!
- Dictionaries are pervasive
- Hash tables implement them efficiently
 - Under an optimistic assumption of random keys
 - Can be “made true” by heuristic hash functions
- Key idea for beating BSTs: Indexing
 - Sacrificed operations: previous, successor
- Chaining strategy for collision resolution
- Next two lectures: speed & space improvements

Summary

- Reduced compare cost to $O(n)/\text{length}$
 - By using a big hash table
 - Or signatures in a small table
- Reduced hash computation to $O(n)/\text{length}$
 - Rolling hash function
- Total cost of phases: $O(n \log n)$
- Not the end: suffix tree achieves $O(n)$

Hashing and Genomes are coming back

Unit	Pset	Week	Date	Lecture (Tuesdays and Thursdays)		Recitation (Wed and Fri)		
Intro Binary Search Trees	PS1 Out: 2/1 Due: Mon 2/14 HW lab: Sun 2/13	1	Tue Feb 01	1	Introduction and Document Distance	1	Python and Asymptotic Complexity	
			Thu Feb 03	2	Peak Finding Problem	2	Peak Finding correctness & analysis	
		2	Tue Feb 08	3	Scheduling and Binary Search Trees	3	Binary Search Tree Operations	
			Thu Feb 10	4	Balanced Binary Search Trees	4	Rotations and AVL tree deletions	
Hashing	PS2 Out: 2/15 Due: Mon 2/28 HW lab: Sun 2/27	3	Tue Feb 15	5	Hashing I : Chaining, Hash Functions	5	Hash recipes, collisions, Python dicts	
			Thu Feb 17	6	Hashing II : Table Doubling, Rolling Hash	6	Probability review, Pattern matching	
		4	Tue Feb 22	-	President's Day - Monday Schedule - No Class	-	No recitation	
			Thu Feb 24	7	Hashing III : Open Addressing	7	Universal Hashing, Perfect Hashing	
Sorting	PS3. Out: 3/1 Due: Mon 3/7 HW lab: Sun 3/6	5	Tue Mar 01	8	Sorting I : Insertion & Merge Sort, Master Theorem	8	Proof of Master Theorem, Examples	
			Thu Mar 03	9	Sorting II : Heaps	9	Heap Operations	
		6	Tue Mar 08	10	Sorting III: Lower Bounds Counting Sort, Radix Sort	10	Models of computation	
			Wed Mar 09	Q1	Quiz 1 in class at 7:30pm. Covers L1-R10. Review Session on Tue 3/8 at 7:30pm.			
Graphs and Search	PS4. Out: 3/10 Due: Fri 3/18 HW lab: W 3/16		Thu Mar 10	11	Searching I: Graph Representation, Depth-1st Search	11	Strongly connected components	
		7	Tue Mar 15	12	Searching II: Breadth-1st Search, Topological Sort	12	Rubik's Cube Solving	
			Thu Mar 17	13	Searching III: Games, Network properties Motifs	13	Subgraph isomorphism	
Shortest Paths	PS5 Out: 3/29 Due: Mon 4/11 HW lab: Sun 4/10	8	Tue Mar 29	14	Shortest Paths I: Introduction, Bellman-Ford	14	Relaxation algorithms	
			Thu Mar 31	15	Shortest Paths II: Bellman-Ford, DAGs	15	Shortest Path applications	
		9	Tue Apr 05	16	Shortest Paths III: Dijkstra	16	Speeding up Dijkstra's algorithm	
			Thu Apr 07	17	Graph applications Genome Assembly	17	Euler Tours	
Dynamic Program ming	PS6 Out: Tue 4/12 Due: Fri 4/29 HW lab: W 4/27	10	Tue Apr 12	18	DP I: Memoization, Fibonacci, Crazy Eights	18	Limits of dynamic programming	
			Wed Apr 13	Q2	Quiz 2 in class at 7:30pm. Covers L11-R17. Review Session on Tue 4/13 at 7:30pm.			
			Thu Apr 14	19	DP II: Shortest Paths, Genome sequence alignment	19	Edit Distance, LCS, cost functions	
Numbers Pictures (NP)	PS7 out Thu 4/28 Due: Fri 5/6 HW lab: Wed 5/4	11	Tue Apr 19	-	Patriot's Day - Monday and Tuesday Off	-	No recitation	
			Thu Apr 21	20	DP III: Text Justification, Knapsack	20	Saving Princess Peach	
		12	Tue Apr 26	21	DP IV: Piano Fingering, Vertex Cover, Structured DP	21	Phylogeny	
			Thu Apr 28	22	Numerics I - Computing on large numbers	22	Models of computation return!	
Beyond		13	Tue May 3	23	Numerics II - Iterative algorithms, Newton's method	23	Computing the nth digit of π	
			Thu May 5	24	Geometry: Line sweep, Convex Hull	24	Closest pair	
			14	Tue May 10	25 Complexity classes, and reductions	25	Undecidability of Life	
			Thu May 12	26	Research Directions (15 mins each) + related classes			
		15	Finals week	Q3	Final exam is cumulative L1-L26. Emphasis on L18-L26. Review Session on Fri 5/13 at 3pm			